

Hadoop Map-Reduce NoSQL Database

By Akash Desai, Nihal Parchand, Viraj Chaudhari

Overview

❖ Challenges:

- Input data is too large (PetaBytes)
- Straightforward repetitive work

❖ Issues:

- Recovering from machine failure
- Debugging and Optimization
- Communication and coordination
- Locality
- Scheduling

Limitations of RDBMS

- ❖ Only structured Data
- ❖ Average sized data (GBs)
- ❖ Scalability (Vertical)
- ❖ Licensed
- ❖ Static schema (Pre-defined)
- ❖ Eg: PostgreSQL, MySQL, Oracle, Microsoft SQL server etc.

Benefits of NoSQL

- ❖ Structured, semi-structured and unstructured data
- ❖ Large datasets (TBs and PBs)
- ❖ Scalability (Horizontal)
- ❖ Dynamic Schema
- ❖ Eg: Cassandra, MongoDB, BigTable, HBase, Neo4j, CouchDB, MapReduce etc.

Introduction of Hadoop

- Apache Open Source Framework
- Stores large heterogeneous dataset on distributed system
- Process data using MapReduce Framework
- Data storage using Hadoop Distributed File System
- Hive Query Language to access data
- Different components - Hive, Pig, MapReduce, Mahout, Latin, Oozie, HDFS etc.

HDFS ARCHITECTURE

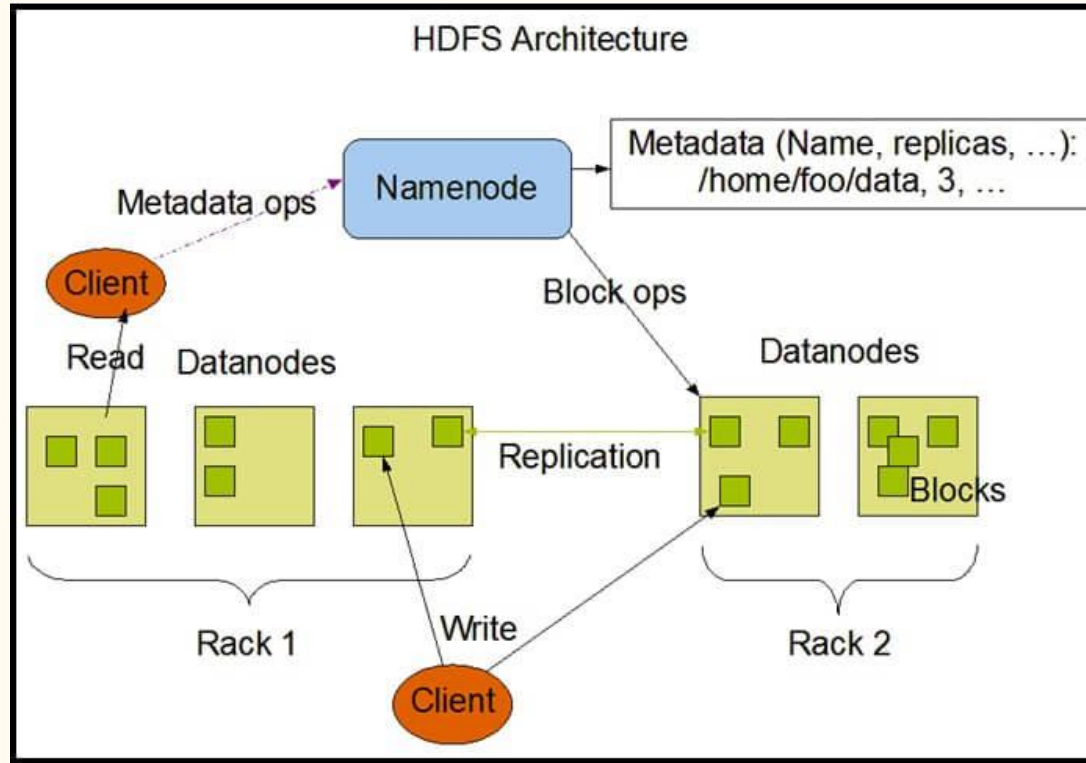


Image source: https://www.mssqltips.com/tipimages2/3180_HDFS_Architecture.jpg

Map-Reduce History

- ❖ Popularized as programming model by Jeffery Dean and Sanjay Ghemawat in 2004.
- ❖ Google used it to collect and analyze website data for search optimizations.
- ❖ Apache started using map-reduce in the subproject Apache Lucern. Provides text search capabilities across large databases.
- ❖ In 2007, Doug Cutting released Hadoop as open source Apache project.

RDBMS

- ❖ Use B-Tree for updates
- ❖ Many read/write operations continually
- ❖ Structured data (XML documents/Tables)
- ❖ Schema-on-write

Map-Reduce

- ❖ Use Sort/Merge for updates
 - ❖ Write once, read many times
 - ❖ Semi-structured data (Spreadsheet) and unstructured data (plain text/image data)
 - ❖ Schema-on-read
-

Map-Reduce Overview

- ❖ Simple programming model (Map + Reduce) that applies to many large-scale computing problems
- ❖ Distributed implementation hides messy details:
 - Fault Tolerance
 - Parallelization
 - I/O Scheduling
 - Network and Disk transfer optimization

Map-Reduce Overview

- ❖ Two types of nodes:
 - Cluster (Same local network and similar hardware)
 - Grid (Shared across geographically and heterogeneous hardware)
- ❖ Data stored in :
 - Filesystem (Unstructured)
 - Database (Structured)

Map-Reduce Operations

- ❖ Map - Process the I/P key-value pair and produce set of O/P intermediate key/value pairs.

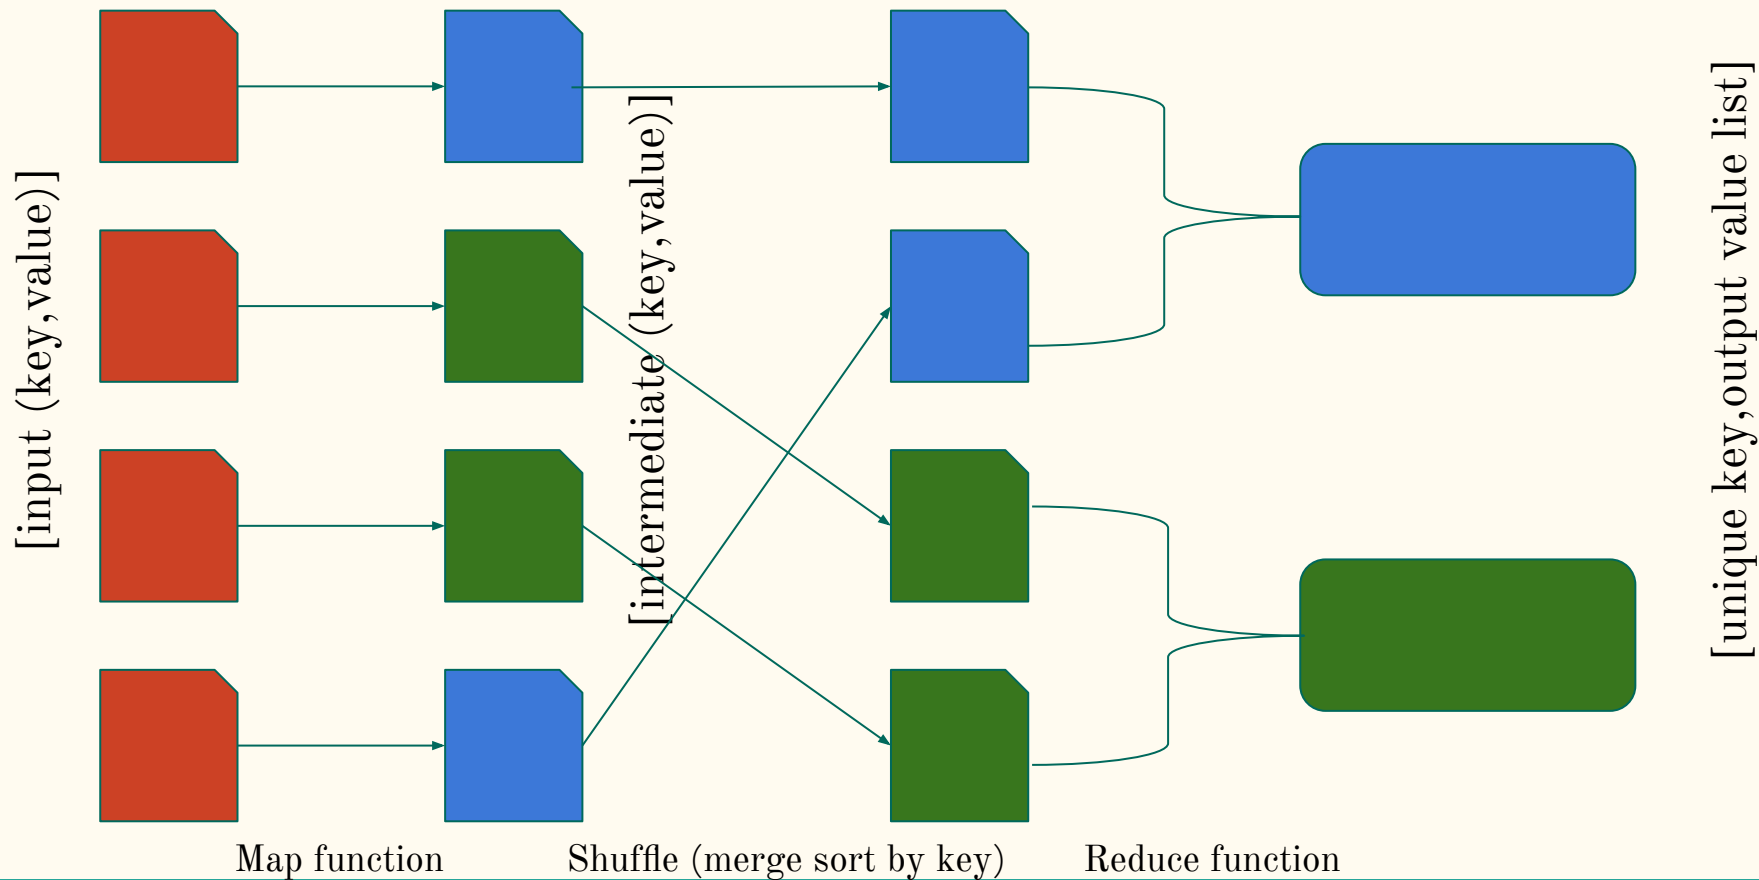
$\text{map}(\text{in_key}, \text{in_value}) \rightarrow \text{list}(\text{out_key}, \text{interm_value})$

- ❖ Shuffle - Use sort/merge to combine intermediate results.

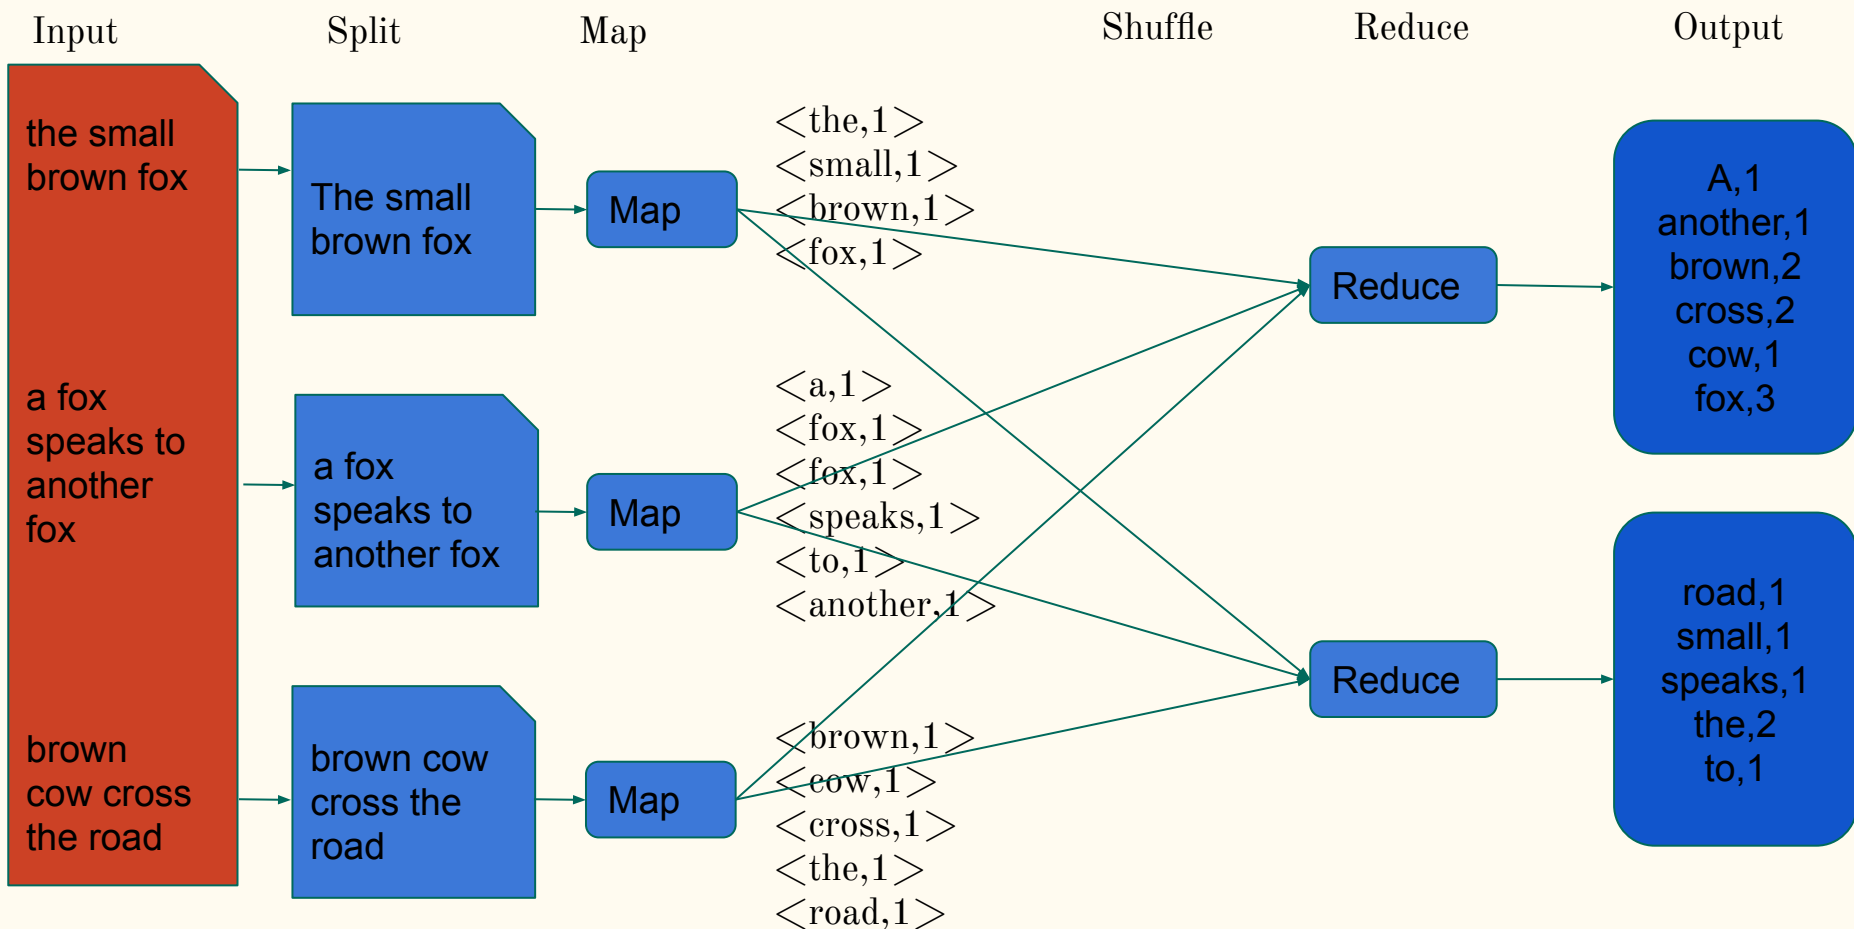
- ❖ Reduce - Process the intermediate key/value, combine intermediate values for each unique key and generate a set of merged output values

$\text{reduce}(\text{out_key}, \text{interm_value}) \rightarrow \text{list}(\text{out_value})$

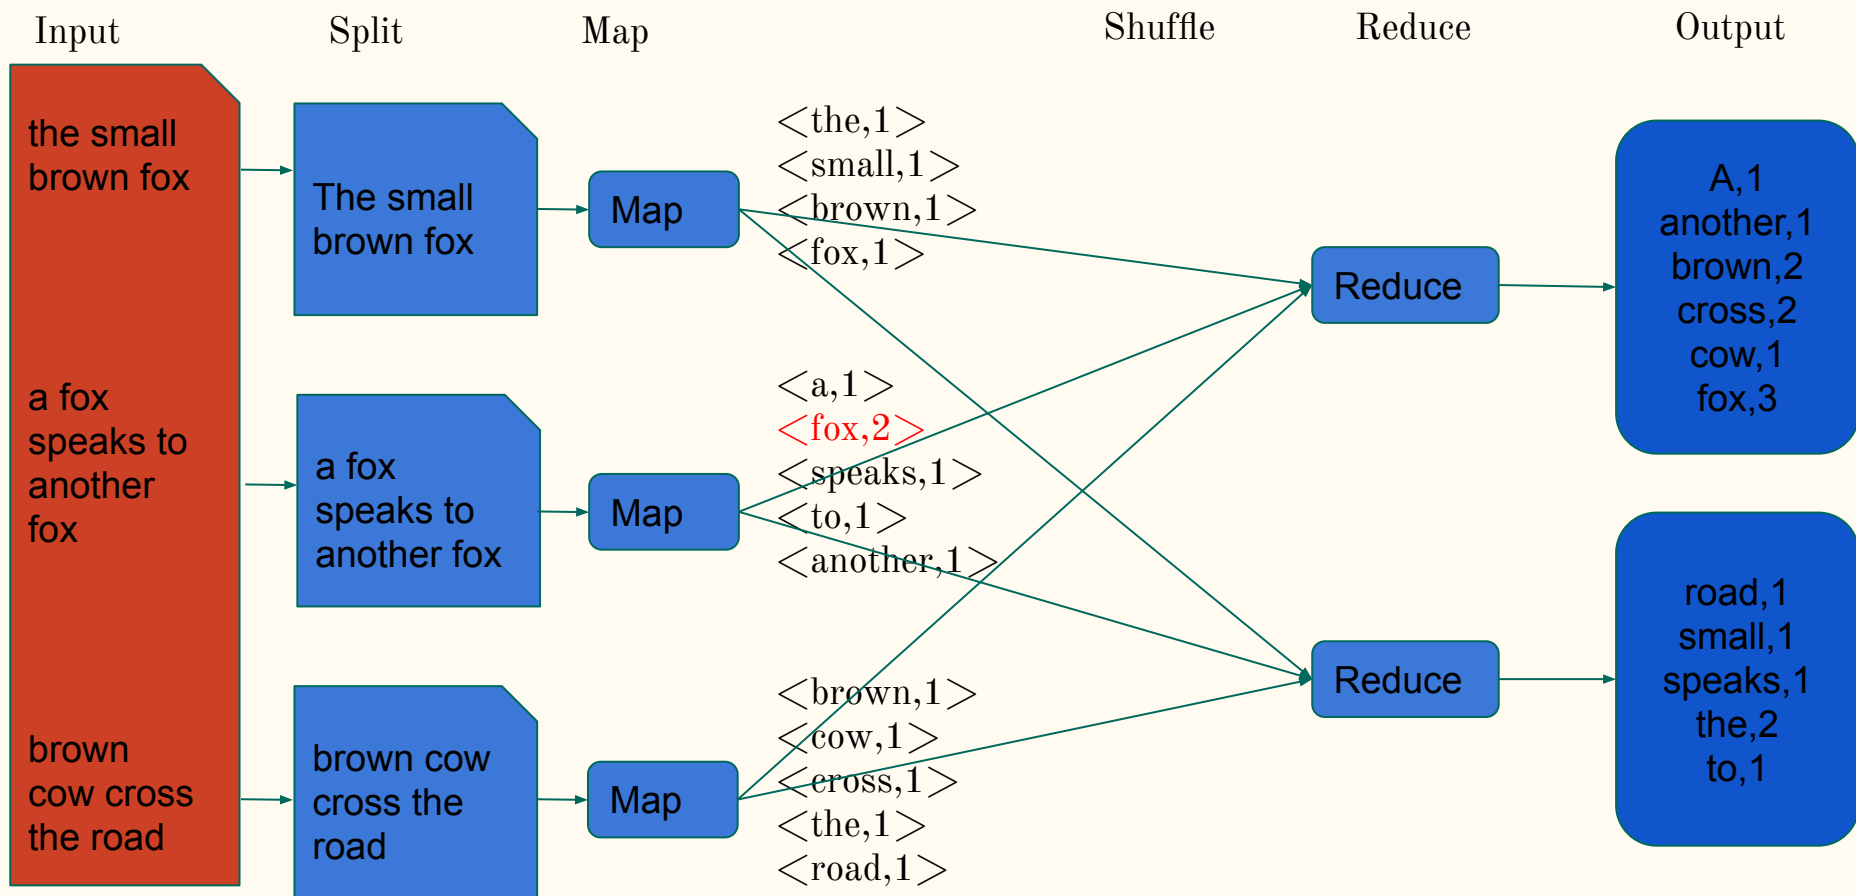
Map-Reduce Programming Model



Map-Reduce WordCount Example



Map-Reduce WordCount Example (Combiner)



Map-Reduce Indexing

- ❖ Index based on File URI
 - Indexed query will be equivalent to full scan query
- ❖ Index based on InputSplit
 - Indexed query better than full scan.
 - Performance is optimized

Map-Reduce Performance

- ❖ Optimized shuffle operation and writing only Map and Reduce function.
- ❖ Combiner - reduce data written to disk
- ❖ Complexity of mapping, shuffle, sorting, and reducing
- ❖ Fault Tolerance can be handled by re-execution
 - Worker failure
 - Master failure

Map-Reduce Refinement

- ❖ Redundant execution
- ❖ Skipping bad records
- ❖ Backup tasks
- ❖ Locality optimization
- ❖ Optional secondary keys for ordering

Transaction Management

ACID

- Atomicity
- Consistency
- Isolation
- Durability

BASE

- Basically available
- Soft state
- Eventual consistency

2 PHASE COMMIT PROTOCOL

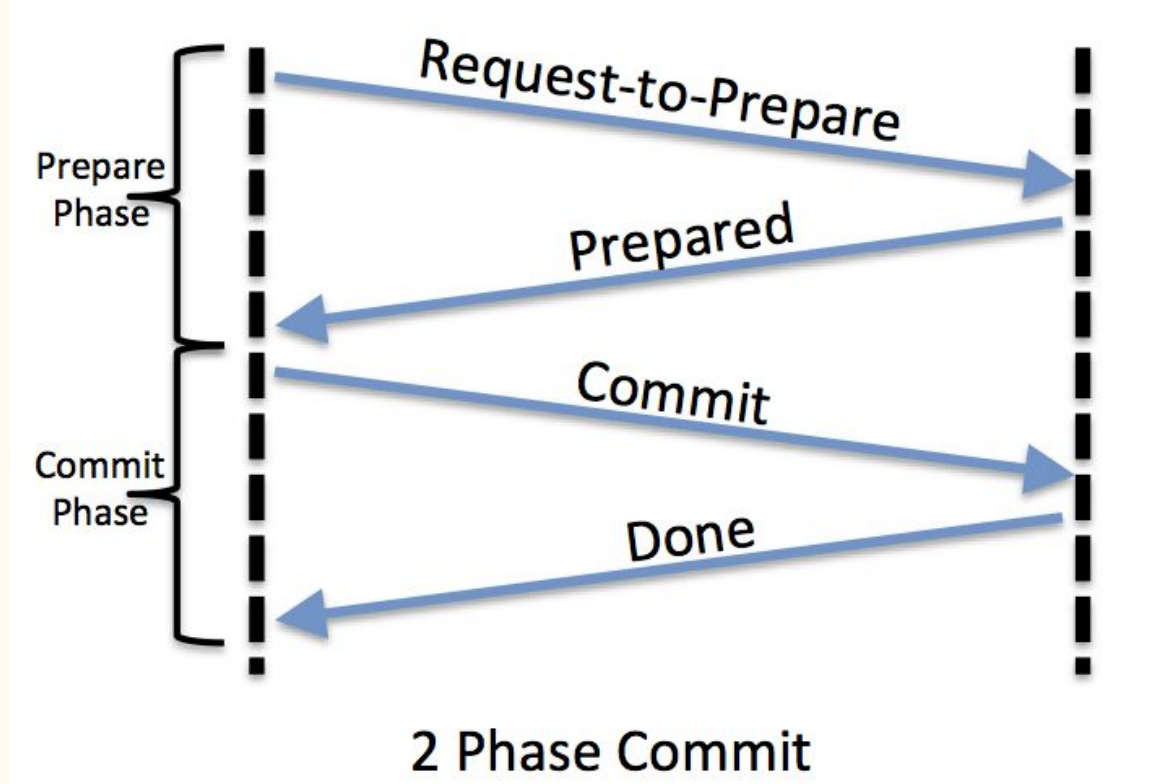


Image source: <https://www.tigerteam.dk/wp-content/uploads/2014/03/2-phase-commit-protocol-flow.png>

Map-Reduce Uses

- ❖ Distributed pattern-based searching
- ❖ Distributed sorting
- ❖ Web link-graph reversal
- ❖ Inverted index construction
- ❖ Document clustering

Application

- Cloudera VM-Oracle Virtual Box virtualization software provides the HDFS instances as well as all the services in the Hadoop framework, which can be installed separately. The components include the HDFS file system, Hive, Impala, MapReduce, PigLatin.
- CentOS operating system provides terminal support.
- City database as an input with columns - ID, Name, CountryCode, District, Population. The database is stored as a CSV input in the HDFS folder.
- We implemented java code, which takes SQL queries as an input and generates output using MapReduce Framework.

RESULT

select name, countrycode, population from city

```
138950 Livonia,USA,100545
138983 Burbank,USA,100316
139018 Clearwater,USA,99936
139052 Midland,USA,98293
139081 Davenport,USA,98256
139111 MissionViejo,USA,98049
139150 MiamiBeach,USA,97855
139184 SunriseManor,USA,95362
139219 NewBedford,USA,94780
139259 ElCajon,USA,94578
139293 Norman,USA,94193
139324 Richmond,USA,94100
139359 Albany,USA,93994
139389 Brockton,USA,93653
139427 Roanoke,USA,93357
139459 Billings,USA,92988
139491 Compton,USA,92864
139525 Gainesville,USA,92291
139560 Fairfield,USA,92256
139596 Arden-Arcade,USA,92040
139635 SanMateo,USA,91799
139670 Visalia,USA,91762
139704 Boulder,USA,91238
139736 Cary,USA,91213
139770 SantaMonica,USA,91084
139808 FallRiver,USA,90555
139847 Kenosha,USA,89447
139880 Elgin,USA,89408
139910 Odessa,USA,89293
139938 Carson,USA,89089
139971 Charleston,USA,89063
140011 CharlotteAmalie,VIR,13000
140051 Harare,ZWE,1410000
```

RESULT

select count(*) from city

```
GC time elapsed (ms)=127
CPU time spent (ms)=2080
Physical memory (bytes) snapshot=534593536
Virtual memory (bytes) snapshot=3147763712
Total committed heap usage (bytes)=479723520
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=140435
File Output Format Counters
  Bytes Written=11
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /output/part-r-00000
```

```
COUNT    4079
```

```
[cloudera@quickstart ~]$ █
```

RESULT

```
select count(*) from city group by countrycode
```

SOM	3
SPM	1
STP	1
SUR	1
SVK	3
SVN	2
SWE	15
SWZ	1
SYC	1
SYR	11
TCA	1
TCD	2
TGO	1
THA	12
TJK	2
TKL	1

RESULT

. select name, countrycode from city where
population < 100000

```
129545 Taikovski,RUS
129575 NovyiUrengoi,RUS
130324 BuonMaThuot,VNM
139018 Clearwater,USA
139052 Midland,USA
139081 Davenport,USA
139111 MissionViejo,USA
139150 MiamiBeach,USA
139184 SunriseManor,USA
139219 NewBedford,USA
139259 ElCajon,USA
139293 Norman,USA
139324 Richmond,USA
139359 Albany,USA
139389 Brockton,USA
139427 Roanoke,USA
139459 Billings,USA
139491 Compton,USA
139525 Gainesville,USA
139560 Fairfield,USA
139596 Arden-Arcade,USA
139635 SanMateo,USA
139670 Visalia,USA
139704 Boulder,USA
139736 Cary,USA
139770 SantaMonica,USA
139808 FallRiver,USA
139847 Kenosha,USA
139880 Elgin,USA
139910 Odessa,USA
139938 Carson,USA
139971 Charleston,USA
140011 CharlotteAmalie,VIR
140408 Rafah,PSE
```

RESULT

select MAX(population) from city

```
DATA_READ=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=140435
File Output Format Counters
  Bytes Written=16
[cloudera@quickstart ~]$ hdfs dfs -cat /output/part-r-00000
RESULT 10500000
[cloudera@quickstart ~]$ █
```

Conclusion

- ❖ Inexpensive (Free to use)
- ❖ Hides all the messy details- Implement only Map and Reduce functions
- ❖ Provides easy model to implement parallel programs
- ❖ Performance of MapReduce is better when the volume of data is large i.e. in PetaBytes.

References

- ❖ MapReduce: Simplified Data Processing on Large Clusters
<https://static.googleusercontent.com/media/research.google.com/en//archive/mapreduce-osdi04.pdf>
- ❖ HADOOP vs RDBMS|Know The 12 Useful Differences
<https://www.educba.com/hadoop-vs-rdbms/>
- ❖ Hadoop: The Definitive Guide
- ❖ <https://en.wikipedia.org/wiki/MapReduce>
- ❖ <https://hadoopi.wordpress.com/2013/05/24/indexing-on-mapreduce-2/>
- ❖ <https://www.slideshare.net/rantav/introduction-to-map-reduce>

Thank You